

HDFS COMMANDS

Command Purpose	Windows Command Example	Mac M1 Equivalent Command Example
-----------------	-------------------------	-----------------------------------

Start Hadoop services	<code>sbin/start-all.cmd</code>	<code>sbin/start-all.sh</code> or <code>sbin/start-dfs.sh &&</code> <code>sbin/start-yarn.sh</code>
Check Hadoop services	<code>jps</code>	<code>jps</code> (same command)
List files/directories	<code>hdfs dfs -ls /path</code>	<code>hdfs dfs -ls /path</code>
Create directory	<code>hdfs dfs -mkdir /path</code>	<code>hdfs dfs -mkdir /path</code>
Create empty file	<code>hdfs dfs -touchz /path</code>	<code>hdfs dfs -touchz /path</code>
Copy from local to HDFS	<code>hdfs dfs -copyFromLocal</code> <code>localfile /path</code> or <code>hdfs dfs</code> <code>-put</code>	<code>hdfs dfs -copyFromLocal localfile</code> <code>/path</code> or <code>hdfs dfs -put localfile</code> <code>/path</code>

Show file contents	<code>hdfs dfs -cat /path</code>	<code>hdfs dfs -cat /path</code>
Copy from HDFS to local	<code>hdfs dfs -copyToLocal /path localdir</code> or <code>hdfs dfs -get /path localdir</code>	<code>hdfs dfs -copyToLocal /path localdir</code> or <code>hdfs dfs -get /path localdir</code>
Show size of files	<code>hdfs dfs -du /dirname</code>	<code>hdfs dfs -du /dirname</code>
Show total size	<code>hdfs dfs -dus /dirname</code>	<code>hdfs dfs -dus /dirname</code>
Remove file	<code>hdfs dfs -rm /path</code>	<code>hdfs dfs -rm /path</code>
Remove directory	<code>hdfs dfs -rmdir /path</code>	<code>hdfs dfs -rmdir /path</code>
Move files within HDFS	<code>hdfs dfs -mv /source /destination</code>	<code>hdfs dfs -mv /source /destination</code>
Copy files within HDFS	<code>hdfs dfs -cp /source /destination</code>	<code>hdfs dfs -cp /source /destination</code>
Append to file	<code>hdfs dfs -appendToFile - /source/destination</code>	<code>hdfs dfs -appendToFile - /source/destination</code>

Count files/directories	<code>hdfs dfs -count /path</code>	<code>hdfs dfs -count /path</code>
Show last part of file	<code>hdfs dfs -tail /data/folder</code>	<code>hdfs dfs -tail /data/folder</code>
Show first part of file	<code>hdfs dfs -head /path</code>	<code>hdfs dfs -head /path</code>
Test if path exists	<code>hdfs dfs -test -e /path</code>	<code>hdfs dfs -test -e /path</code>
Test if directory	<code>hdfs dfs -test -d /path</code>	<code>hdfs dfs -test -d /path</code>
Test if file	<code>hdfs dfs -test -f /path</code>	<code>hdfs dfs -test -f /path</code>
Test if file empty	<code>hdfs dfs -test -z /path</code>	<code>hdfs dfs -test -z /path</code>
Move from local to HDFS	<code>hdfs dfs -moveFromLocal local_source /hdfs_destination</code>	<code>hdfs dfs -moveFromLocal local_source /hdfs_destination</code>
Merge files	<code>hdfs dfs -getmerge -nl /source /localdestination</code>	<code>hdfs dfs -getmerge -nl /source /localdestination</code>

Get checksum	<code>hdfs dfs -checksum /file</code>	<code>hdfs dfs -checksum /file</code>
Change group	<code>hdfs dfs -chgrp new_group_name /file</code>	<code>hdfs dfs -chgrp new_group_name /file</code>
Show last modified time	<code>hdfs dfs -stat /path</code>	<code>hdfs dfs -stat /path</code>
Empty trash	<code>hdfs dfs -expunge</code>	<code>hdfs dfs -expunge</code>
Change owner/group	<code>hdfs dfs -chown owner: group /path</code>	<code>hdfs dfs -chown owner: group /path</code>
Change permissions	<code>hdfs dfs -chmod 777 /path</code>	<code>hdfs dfs -chmod 777 /path</code>
Set replication factor	<code>hdfs dfs -setrep -w n /path</code>	<code>hdfs dfs -setrep -w n /path</code>

- **Basic HDFS Commands:**

- **jps:** Lists running Hadoop Java processes.
- **ls:** Lists files and directories in HDFS.
Syntax: `hdfs dfs -ls /path`
- **mkdir:** Creates a directory in HDFS.
Syntax: `hdfs dfs -mkdir /path`
- **touchz:** Creates an empty file in HDFS.
Syntax: `hdfs dfs -touchz /path`

- **copyFromLocal / put:** Copies files from local filesystem to HDFS.
Syntax: `hdfs dfs -copyFromLocal /localpath /hdfspath` or `hdfs dfs -put /localpath /hdfspath`
- **cat:** Displays file content.
Syntax: `hdfs dfs -cat /path`
- **copyToLocal / get:** Copies files from HDFS to local filesystem.
Syntax: `hdfs dfs -get /hdfspath /localpath`
- **du:** Shows size of files/directories.
Syntax: `hdfs dfs -du /dirname`
- **rm:** Removes a file from HDFS.
Syntax: `hdfs dfs -rm /path`
- **rmdir:** Removes a directory in HDFS.
Syntax: `hdfs dfs -rmdir /path`
- **mv:** Moves or renames files/directories within HDFS.
Syntax: `hdfs dfs -mv /source /destination`
- **cp:** Copies files/directories within HDFS.
Syntax: `hdfs dfs -cp /source /destination`
- **appendToFile:** Appends data to an existing HDFS file.
Syntax: `hdfs dfs -appendToFile - /path`
- **count:** Counts files, directories, and bytes under a path.
Syntax: `hdfs dfs -count /path`
- **tail:** Shows the last part of a file.
Syntax: `hdfs dfs -tail /path`
- **head:** Shows the first part of a file.
Syntax: `hdfs dfs -head /path`
- **test:** Checks file/directory status.
 - `-e` tests if path exists
 - `-d` tests if path is a directory
 - `-f` tests if path is a file
 - `-z` tests if file is empty (zero bytes)
 Syntax: `hdfs dfs -test -e /path`
- **moveFromLocal:** Moves files from local filesystem to HDFS (deletes local after move).
Syntax: `hdfs dfs -moveFromLocal /local_source /hdfs_destination`
- **getmerge:** Merges multiple files in HDFS into a single local file.
Syntax: `hdfs dfs -getmerge -nl /source /localdestination`
- **checksum:** Retrieves checksum of a file to verify data integrity.
Syntax: `hdfs dfs -checksum /file`
- **chgrp:** Changes group ownership of files/directories.
Syntax: `hdfs dfs -chgrp new_group /file`

- **stat:** Shows last modification time of a file/directory.
Syntax: `hdfs dfs -stat /path`
- **expunge:** Empties the HDFS trash, permanently deleting files.
Syntax: `hdfs dfs -expunge`
- **chown:** Changes owner and group of files/directories.
Syntax: `hdfs dfs -chown owner:group /path`
- **chmod:** Changes permissions of files/directories.
Syntax: `hdfs dfs -chmod 777 /path`
(Permissions: Read=4, Write=2, Execute=1)
- **setrep:** Sets replication factor for files/directories.
Syntax: `hdfs dfs -setrep -w <n> /path`
(**-w** waits for replication to complete, **n** is replication count)

HIVE COMMANDS

- Hive is a data warehousing infrastructure built on Hadoop, enabling SQL-like querying (HiveQL) over large datasets stored in HDFS.
- It converts HiveQL queries into MapReduce jobs for batch processing of big data.
- Hive is designed for batch processing, not real-time or OLTP workloads.

HiveQL Command Categories:

- DDL (Data Definition Language): CREATE, DROP, ALTER, TRUNCATE
- DML (Data Manipulation Language): INSERT, SELECT, UPDATE, DELETE
- Partitioning: Efficiently manages large datasets by dividing tables into partitions.
- Joins: Supports INNER, LEFT, RIGHT, FULL OUTER joins.
- Aggregation: COUNT, SUM, AVG, MIN, MAX functions.
- Sorting & Filtering: ORDER BY, WHERE, HAVING clauses.

Database Operations:

- Create database: `CREATE DATABASE ecommerce_db;`
- Use database: `USE ecommerce_db;`
- Show databases: `SHOW DATABASES;`
- Drop database: `DROP DATABASE ecommerce_db CASCADE;`

Table Operations:

- Create table example:
- sql

```
CREATE TABLE products (  
  product_id INT,  
  name STRING,  
  category STRING,  
  price FLOAT  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
STORED AS TEXTFILE;
```

-
- Show tables: `SHOW TABLES;`
- Describe table: `DESCRIBE products;`
- Drop table: `DROP TABLE products;`

Data Operations:

- Insert data into table:
- sql

```
INSERT INTO TABLE products VALUES (1, 'Laptop', 'Electronics', 800.50), (2, 'Smartphone', 'Electronics', 500.75);
```

-
- Load data from HDFS:
- sql

```
LOAD DATA INPATH '/user/hive/products.csv' INTO TABLE products;
```

-
- Select data: `SELECT * FROM products;`
- Filter data: `SELECT * FROM products WHERE category = 'Electronics';`

Joins in HiveQL:

- Inner Join:
- sql

```
SELECT p.name, s.amount FROM products p JOIN sales s ON p.product_id = s.product_id;
```

-
- Left Join, Right Join, Full Outer Join supported similarly.

Aggregation and Sorting:

- Group By example:
- sql

```
SELECT category, COUNT(*) AS product_count FROM products GROUP BY category;
```

-
- Order By example:
- sql

```
SELECT * FROM products ORDER BY price DESC;
```

-
- Limit example: `SELECT * FROM products LIMIT 5;`

Views:

- Create view:
- sql


```
CREATE VIEW electronics_view AS SELECT * FROM products WHERE category = 'Electronics';
```

-
- Use and drop views supported.

Operators in HiveQL:

- Arithmetic: +, -, *, /, %
- Comparison: =, !=, >, <, >=, <=
- Logical: AND, OR, NOT
- String: LIKE, REGEXP, IN

Complex Data Types:

- Arrays, Maps, Structs supported for complex data modeling.
- Example:
- sql

```
CREATE TABLE temperature (  
  sno INT,  
  place STRING,  
  mytemp ARRAY<DOUBLE>  
)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY '\t'
```

- COLLECTION ITEMS TERMINATED BY ',';

Hive Data Types:

- *Primitive Types*: TINYINT (1 byte), SMALLINT (2 bytes), INT (4 bytes), FLOAT (4 bytes), DOUBLE (8 bytes)
- *Date/Time Types*: TIMESTAMP (with nanosecond precision), DATE (YYYY-MM-DD format)
- *String Types*: STRING (sequence of characters), VARCHAR (1 to 65535 length), CHAR (fixed length up to 255)
- *Complex Types*: Arrays (homogeneous collections), Maps (key-value pairs with primitive keys), Structs (nested complex structures)

Hive Bucketing:

- Improves query performance by hashing data into buckets using **CLUSTERED BY** clause.
- Buckets determine the number of files and ensure even data distribution.
- Often used together with partitioning for faster queries.

- Example to create bucketed table:
- sql

```
CREATE TABLE table_name (column1 DATA_TYPE, column2 DATA_TYPE)
CLUSTERED BY (column_name) INTO N BUCKETS
STORED AS FILE_FORMAT;
```

- Enable bucketing with `SET hive.enforce.bucketing = true;`

Hive Partitioning:

- Splits data into partitions based on a partition column to reduce scan times and improve query performance.
- Supports static and dynamic partitioning.
- Example to create partitioned table:
- sql

```
CREATE TABLE table_name (column1 DATA_TYPE, column2 DATA_TYPE)
PARTITIONED BY (partition_column DATA_TYPE)
STORED AS FILE_FORMAT;
```

- Insert data into partition:
- sql

```
INSERT INTO table_name PARTITION (partition_column=value)
VALUES (value1, value2);
```

-

Basic Hive Operations:

- Create database: `CREATE DATABASE IF NOT EXISTS abc;`
- Show databases: `SHOW DATABASES;`
- Create table with delimiter and storage format:
- sql

```
CREATE TABLE customers(id INT, fname STRING, lname STRING,
city STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '|'
STORED AS TEXTFILE;
```

- Load data into table:
- sql

```
LOAD DATA LOCAL INPATH 'path_to_file' INTO TABLE customers;
```

- Query data: `SELECT * FROM table;`
- Drop table: `DROP TABLE table_name;`
- Rename table: `ALTER TABLE old_name RENAME TO new_name;`
- Add columns: `ALTER TABLE table_name ADD COLUMNS (salary INT);`

HBASE COMMANDS

Apache HBase is an open-source, distributed, scalable, column-oriented NoSQL database built on top of the Hadoop Distributed File System (HDFS). It is modeled after Google's Bigtable and designed to provide random, real-time read/write access to very large datasets, often spanning billions of rows and millions of columns.

Key Features:

- Linear and modular scalability to handle massive data growth.
- Strictly consistent reads and writes for reliable data operations.
- Automatic sharding (region splitting) and load balancing across cluster nodes.
- High availability with automatic failover support for RegionServers.
- Integration with Hadoop ecosystem and MapReduce for batch processing.
- Provides easy-to-use Java APIs and supports REST and Thrift gateways.
- Supports block caching and Bloom filters for fast query performance.

Architecture Components:

- HMaster: Manages cluster coordination, schema changes, region assignments, and load balancing.
- RegionServers: Store and manage data regions, handle client read/write requests, and perform background compactions.
- ZooKeeper: Coordinates distributed cluster state, leader election, and failover processes.
- HDFS: Underlying storage layer providing fault tolerance and replication.

Use Cases:

- Suitable for applications requiring fast random access to large-scale structured or semi-structured data.
- Used by companies like Facebook, Twitter, Yahoo, and Adobe for heavy write and real-time data access workloads.

Limitations:

- Requires careful upfront schema design; lacks ad-hoc query flexibility.
- Does not support joins or full ACID transactions across multiple regions.
- Not optimized for small datasets or complex relational queries.

Comparison with HDFS:

- HDFS is a distributed file system optimized for high-throughput batch processing and sequential data access.
- HBase provides low-latency random read/write access on top of HDFS, enabling real-time querying of big data.

Creating tables and adding data to it:

- a. create 'person','data' //create table
- b. put 'person',1,'data:name','Sahil' //put data into table
- c. put 'person',1,'data:city','washington'
- d. put 'person',1,'data:id','10'
- e. scan 'person' // to print the data of person table
- f. get 'person','1' //get data from a particular row
- g. get 'person','1',{COLUMN => 'data:id'} //fetch data of particular column in a column family
- h. get 'person','1',{COLUMN => ['data:id','data:name']}
//fetch data from multiple columns in a column family

Alter table:

- a. alter 'table1', {NAME => 'COLFAM2'}//Adding column family
- b. alter 'table1',{NAME => 'COLFAM2' , METHOD => 'delete'} // DELETING A COLUMN FAMILY
- c. alter 'table1',{NAME => 'COLFAM2' , VERSIONS => 2}//CHANGING THE VERSIONS OF THE COLUMN FAMILY
- d. alter 'student', READONLY //for making table read-only
- e. alter 'student', {NAME => 'semsester', VERSIONS => 5}
- f. alter 'student', MAXFILESIZE='65165'
- g. alter_status 'student' ///how many regions of the table have been altered
- h. create_namespace 'mynamespace' //creates a namespace(logical grouping) of the tables in hbase
- i. create 'mynamespace:table_one', 'colfam1' //create a table inside a namespace
- j. describe_namespace 'mynamespace' //describe namespace
- k. list_namespace //list all namespaces

PIG COMMANDS

Apache Pig is a high-level platform and scripting language designed to analyze and process large datasets on Hadoop. It provides a language called Pig Latin, which simplifies writing data analysis programs by abstracting complex MapReduce jobs into easier-to-write scripts.

Key Features:

- **Ease of Programming:** Pig Latin is a procedural, SQL-like language that enables programmers to express data transformations as data flows, making it easier to write, understand, and maintain complex data processing tasks.
- **Abstraction over MapReduce:** Pig scripts are automatically compiled into a series of MapReduce jobs, allowing users to focus on data logic rather than low-level programming.
- **Optimization:** The Pig framework optimizes the execution plan automatically to improve performance.
- **Extensibility:** Users can write their own functions (UDFs) for custom processing.
- **Supports Semi-structured Data:** Pig can handle both structured and semi-structured data, supporting complex data types like tuples, bags, and maps.

Architecture Components:

- **Parser:** Parses Pig Latin scripts and generates a logical plan represented as a directed acyclic graph (DAG).
- **Optimizer:** Optimizes the logical plan to reduce resource usage and execution time.
- **Compiler:** Converts the optimized plan into MapReduce jobs.
- **Execution Engine:** Executes the MapReduce jobs on a Hadoop cluster and returns results.

Use Cases:

- **Ideal for data transformation, ETL (Extract, Transform, Load) tasks, and prototyping MapReduce jobs without writing Java code.**
- **Useful when working with large volumes of data that require parallel processing.**

Comparison with Hive:

- **Pig is procedural and better suited for data pipelines and transformations.**

- Hive uses a declarative SQL-like language (HiveQL) and is preferred for ad-hoc querying and reporting on structured data.
- Pig generally offers more control over data flow and can be faster for certain operations.

Starting Apache Pig

- Command (Windows & Mac):
- `bash`

```
pig -x local
```

-
- Opens the Grunt shell for interactive Pig Latin commands.

Basic Pig Latin Commands

- Load Data from Local File:
- `text`

```
data = LOAD 'C:/Users/gurvi/Downloads/student.txt' USING
PigStorage(',') AS (id:int, name:chararray, marks:int);
```

-
- (On Mac, use Unix-style path, e.g., `/Users/gurvi/Downloads/student.txt`)
- Display Data:
- `text`

```
DUMP data;
```

-
- Store Data to Local Directory:
- `text`

```
STORE data INTO 'C:/Users/gurvi/Documents/my_output' USING
PigStorage(',');
```

-
- (On Mac, use Unix-style path)

Data Filtering

- Filter rows based on condition:

- text

```
high_scorers = FILTER data BY marks > 70;
DUMP high_scorers;
```

-
- Filter with multiple conditions:
- text

```
mid_range = FILTER data BY marks >= 60 AND marks <= 80;
DUMP mid_range;
```

-
- Filter by string equality:
- text

```
john_data = FILTER data BY name == 'John';
DUMP john_data;
```

-

Sorting Data

- Sort ascending:
- text

```
sorted_data = ORDER data BY marks ASC;
DUMP sorted_data;
```

-
- Sort descending:
- text

```
sorted_data_desc = ORDER data BY marks DESC;
DUMP sorted_data_desc;
```

-

Handling NULL Values

- Remove rows where marks is NULL:
- text

```
valid_data = FILTER data BY marks IS NOT NULL;
DUMP valid_data;
```

- Remove rows where name or marks is NULL:

- text

```
clean_data = FILTER data BY name IS NOT NULL AND marks IS NOT NULL;
```

-
- Replace NULL values:
- text

```
cleaned = FOREACH data GENERATE
    id,
    (name IS NOT NULL ? name : 'Unknown') AS name,
    (marks IS NOT NULL ? marks : 0) AS marks;
DUMP cleaned;
```

-

Grouping and Aggregation

- Group data by marks:
- text

```
grouped_data = GROUP data BY marks;
```

-
- Count students per marks group:
- text

```
counted = FOREACH grouped_data GENERATE group AS marks,
COUNT(data) AS student_count;
DUMP counted;
```

-

Joining Data

- Load two datasets:
- text

```
students = LOAD 'C:/Users/gurvi/Downloads/students.txt' USING
PigStorage(',') AS (id:int, name:chararray);
scores = LOAD 'C:/Users/gurvi/Downloads/scores.txt' USING
PigStorage(',') AS (id:int, marks:int);
```

-

- Join on student ID:
- text

```
joined_data = JOIN students BY id, scores BY id;  
final = FOREACH joined_data GENERATE students::id, name,  
marks;  
DUMP final;
```

-

Field Access and Type Casting

- Access first field:
- text

```
ids = FOREACH data GENERATE $0;
```

-

- Cast fields to specific types:
- text

```
converted = FOREACH data GENERATE  
  (int)$0 AS id,  
  (chararray)$1 AS name,  
  (int)$2 AS marks;
```