

CA -2

Andol Saketh

12221210

Roll no: - 35

Q1:

```
Administrator: Command Prompt - hbase shell
at org.jruby.Ruby.runScript(Ruby.java:692)
at org.jruby.Ruby.runScript(Ruby.java:690)
at org.jruby.Ruby.runNormally(Ruby.java:597)
at org.jruby.Ruby.runFromMain(Ruby.java:446)
at org.jruby.Main.doRunFromMain(Main.java:369)
at org.jruby.Main.internalRun(Main.java:258)
at org.jruby.Main.run(Main.java:224)
at org.jruby.Main.run(Main.java:208)
at org.jruby.Main.main(Main.java:188)

ERROR: Failed 1 action: Table 'student' was not found, got: secondtable.: 1 time, servers with issues: null

Here is some help for this command:
Put a cell 'value' at specified table/row/column and optionally
timestamp coordinates. To put a cell value into table 'ns1:t1' or 't1'
at row 'r1' under column 'c1' marked with the time 'ts1', do:

hbase> put 'ns1:t1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value'
hbase> put 't1', 'r1', 'c1', 'value', ts1
hbase> put 't1', 'r1', 'c1', 'value', (ATTRIBUTES=>{'mykey'=>'myvalue'})
hbase> put 't1', 'r1', 'c1', 'value', ts1, (ATTRIBUTES=>{'mykey'=>'myvalue'})
hbase> put 't1', 'r1', 'c1', 'value', ts1, (VISIBILITY=>'PRIVATE|SECRET')

The same commands also can be run on a table reference. Suppose you had a reference
t to table 't1', the corresponding command would be:

hbase> t.put 'r1', 'c1', 'value', ts1, (ATTRIBUTES=>{'mykey'=>'myvalue'})

hbase(main):007:0> put 'Library', '1','Books','student'
0 row(s) in 0.0130 seconds
```

```
Administrator: Command Prompt - hbase shell
Microsoft Windows [Version 10.0.26100.3915]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd C:\hadoopsetup\hadoop-3.2.4\sbin

C:\hadoopsetup\hadoop-3.2.4\sbin>start-all.cmd
This script is deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoopsetup\hadoop-3.2.4\sbin>cd "C:\hbase\hbase-1.4.9-bin\hbase-1.4.9\bin"

C:\hbase\hbase-1.4.9-bin\hbase-1.4.9\bin>start-hbase.cmd

C:\hbase\hbase-1.4.9-bin\hbase-1.4.9\bin>hbase shell
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:C:/hbase\hbase-1.4.9-bin\hbase-1.4.9\lib\slf4j-log4j12-1.7.10.jar/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:C:/hadoop\hadoop-3.2.4\share\hadoop\common\lib\slf4j-reload4j-1.7.35.jar/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.9, r6625b212e46d01cb17db9ac2e9e927fdb201afa1, Wed Dec 5 11:54:10 PST 2018

hbase(main):001:0> create 'Library','Books'
0 row(s) in 1.5050 seconds

=> hbase::Table - Library
hbase(main):002:0> put 'Library','1','Books','student'
0 row(s) in 0.0700 seconds

hbase(main):003:0> alter 'Library',{NAME => 'Books',VERSIONS => 4}
SyntaxError: (hbase):3: syntax error, unexpected tLCURLY

alter 'Library',{NAME => 'Books',VERSIONS => 4}
^

hbase(main):004:0> alter 'Library',{NAME => 'Books',VERSIONS => 4}
updating all regions with the new schema...
0/1 regions updated.
1/1 regions updated.
done.
0 row(s) in 3.2710 seconds

hbase(main):005:0> put 'Library',1,'Books','book_title'
0 row(s) in 0.0290 seconds

hbase(main):006:0> put 'Library',1,'student','student_id'
```

Q2: -

```
Administrator: Command Prompt - hive
2025-04-30T14:52:30,016 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> LOAD DATA LOCAL INPATH 'C:\Users\ayush\Downloads\course_info.txt' INTO TABLE employee_info PARTITION (course_name='Btech');
2025-04-30T14:53:39,018 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:53:38,018 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
FAILED: ValidationFailureSemanticException course_info table is not partitioned but partition spec exists: (course_name=Btech)
2025-04-30T14:53:38,045 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:53:38,045 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> LOAD DATA LOCAL INPATH 'C:\Users\ayush\Downloads\course_info.txt' INTO TABLE student_data PARTITION (course_name='Btech');
2025-04-30T14:54:46,670 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:54:46,670 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
loading data to table course_info.student_data partition (course_name=Btech)
2025-04-30T14:54:46,820 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.common.FileUtils - Creating directory if it doesn't exist: hdfs://localhost:9820/user/hive/warehouse/course_info.db/student_data/course_name=Btech
OK
Time taken: 0.432 seconds
2025-04-30T14:54:47,104 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:54:47,104 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> SET hive.exec.dynamic.partition = true;
2025-04-30T14:55:24,470 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:55:24,482 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
2025-04-30T14:55:24,497 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:55:24,497 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> SET hive.exec.dynamic.partition.mode = nonstrict;
2025-04-30T14:55:40,841 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:55:40,842 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
2025-04-30T14:55:40,840 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:55:40,844 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> SHOW PARTITIONS student_data;
2025-04-30T14:56:02,507 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:56:02,507 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
OK
course_name=Btech
Time taken: 0.113 seconds, Fetched: 1 row(s)
2025-04-30T14:56:02,628 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:56:02,629 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.q1.session.SessionState - Resetting thread name to main
hive> SELECT * FROM student_data WHERE course_name = 'Btech';
2025-04-30T14:56:36,811 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: d853a9d9-2ebd-400b-ac99-7f00d41867ef
2025-04-30T14:56:36,812 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to d853a9d9-2ebd-400b-ac99-7f00d41867ef main
2025-04-30T14:56:37,160 INFO [d853a9d9-2ebd-400b-ac99-7f00d41867ef main] org.apache.hadoop.hive.common.FileUtils - Creating directory if it doesn't exist: hdfs://localhost:9820/tmp/hive/ayush/d853a9d9-2ebd-400b-ac99-7f00d41867ef/hive_2025-04-30_14-56-36_813_5338465977942971605-1/-mr-10801/hive-staging_hive_2025-04-30_14-56-36_813_5338465977942971605-1
OK
Btech NULL Btech
BBA NULL Btech
Btech NULL Btech
Btech NULL Btech
Btech NULL Btech
Btech NULL Btech
BBA NULL Btech
Btech NULL Btech
Btech NULL Btech
Time taken: 0.393 seconds, Fetched: 10 row(s)
```

Q3: -

```
Administrator: Command Prompt - hive
Microsoft Windows [Version 10.0.26100.3915]
(c) Microsoft Corporation. All rights reserved.

C:\Windows\System32>cd C:\hadoopsetup\hadoop-3.2.4\sbin

C:\hadoopsetup\hadoop-3.2.4\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoopsetup\hadoop-3.2.4\sbin>cd C:\hive\apache-hive-3.1.3-bin\apache-hive-3.1.3-bin\bin

C:\hive\apache-hive-3.1.3-bin\apache-hive-3.1.3-bin\bin>hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar-file:C:/hive/apache-hive-3.1.3-bin/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar-file:C:/hbase/hbase-1.4.9-bin/hbase-1.4.9/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar-file:C:/hadoopsetup/hadoop-3.2.4/share/hadoop/common/lib/slf4j-reload4j-1.7.35.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2025-04-30T15:21:34,930 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Found configuration file file:/C:/hive/apache-hive-3.1.3-bin/conf/hive-site.xml
Hive Session ID = 48f11955-be9c-41f7-8004-6ed5443f69d5

Logging initialized using configuration in jar-file:C:/hive/apache-hive-3.1.3-bin/lib/hive-common-3.1.3.jar!/hive-log4j2.properties Async: true
2025-04-30T15:21:37,500 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Created HDFS directory: /tmp/hive/saket/48f11955-be9c-41f7-8004-6ed5443f69d5
2025-04-30T15:21:37,513 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Created local directory: C:/Users/saket/AppData/Local/Temp/saket/48f11955-be9c-41f7-8004-6ed5443f69d5
2025-04-30T15:21:37,519 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Created HDFS directory: /tmp/hive/saket/48f11955-be9c-41f7-8004-6ed5443f69d5/_tmp_space.db
2025-04-30T15:21:37,523 INFO [main] org.apache.hadoop.hive.conf.HiveConf - Using the default value passed in for log id: 48f11955-be9c-41f7-8004-6ed5443f69d5
2025-04-30T15:21:37,523 INFO [main] org.apache.hadoop.hive.q1.session.SessionState - Updating thread name to 48f11955-be9c-41f7-8004-6ed5443f69d5 main
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Hive Session ID = a99af4c3-49e3-41d4-842f-74074b951373
2025-04-30T15:21:46,690 INFO [48f11955-be9c-41f7-8004-6ed5443f69d5 main] CLIDriver - Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
2025-04-30T15:21:46,710 INFO [pool-9-thread-1] org.apache.hadoop.hive.q1.session.SessionState - Created HDFS directory: /tmp/hive/saket/a99af4c3-49e3-41d4-842f-74074b951373
2025-04-30T15:21:46,715 INFO [pool-9-thread-1] org.apache.hadoop.hive.q1.session.SessionState - Created local directory: C:/Users/saket/AppData/Local/Temp/saket/a99af4c3-49e3-41d4-842f-74074b951373
2025-04-30T15:21:46,722 INFO [pool-9-thread-1] org.apache.hadoop.hive.q1.session.SessionState - Created HDFS directory: /tmp/hive/saket/a99af4c3-49e3-41d4-842f-74074b951373/_tmp_space.db
hive> create employee_data(emp_id INT, emp_name STRING, department STRING, salary INT, joining_date INT, manager_id INT, bonus INT)
> > describe employee_data
> > Insert(101,Alice,HR,55000,2022-01-15,201,5000)
> > Insert(102,Bob,IT,70000,2021-11-20,202,6000)
> > Insert(103,Charlie,IT,80000,2020-07-30,203,10000)
> > Insert(104,David,Sales,45000,2023-02-10,204,2000)
> > Insert(104,David,Sales,45000,2023-02-10,204,2000)
```